

# **Concentrated Differentially Private Gradient Descent with Adaptive per-Iteration Privacy Budget**

**Jaewoo Lee & Daniel Kifer**

**KDD 2018**

**Presented by Christian McDaniel**

## **Introduction**

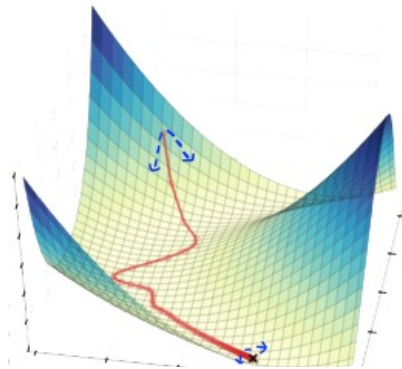
# Introduction

| Iterative optimization algorithm  | Differentially Private iterative algorithms   |
|---|---|
| <p>+ min <math>f</math> by finding optimal parameter vector <math>w^*</math></p> <p>1) Initialize <math>w_0</math><br/> 2) Generate <math>\{w_t\}</math> s.t. <math>w_t</math> tends toward optimal <math>w</math> as <math>t \rightarrow \infty</math></p> <p>e.g., Gradient Descent</p> $w_{t+1} = w_t - \alpha_t(\nabla f(w_t))$ | <p>Gradient Descent under Differential Privacy</p> $w_{t+1} = w_t - \alpha_t(\nabla f(w_t) + Y_t)$ <p>Prior efforts: Naive conversion</p> <p>1) Pick <math>T</math> a priori<br/> 2) Split the budget evenly</p> $\epsilon_1 = \dots = \epsilon_T = \epsilon/T.$ <p>3) Update parameters using a noisy gradient</p> |

# Introduction

## Naive Privacy conversion has a few issues:

- Accuracy depends on  $T$ 
  - Too Small v. Too Large
- The nature of the gradient changes as the optimum is approached
  - Early v. Near Optimum



# Introduction

## Naive Privacy conversion has a few issues:

- Accuracy depends on  $T$ 
  - Too Small v. Too Large
- The nature of the gradient changes as the optimum is approached
  - Early v. Near

## Authors Propose:

- (First known) Adaptive Privacy Budget for zero-mean Concentrated Differential Privacy (more later)
  - re: Larger v. Smaller norm gradients
  - Allows an adaptive  $T$

# Introduction

**Naive Privacy conversion has a few issues:**

- Accuracy depends on  $T$ 
  - Too Small v Too Large
- The nature of the gradient changes as the optimum is approached
  - Early v. Near

**Authors Propose:**

- Adaptive Privacy Budget for zCDP
  - At each iteration:
    - Partial allocation to calculating the noisy gradient:

$$\tilde{S}_t = \nabla f(\mathbf{w}_t) + Y_t$$

## Authors Propose:

- Adaptive Privacy Budget for zCDP
  - At each iteration:
    - Partial allocation to calculating the noisy gradient
    - Allocate the rest to determining the step size:
      - Minimize

$$f(\mathbf{w}_t - \alpha \tilde{S}_t)$$

- Update accordingly:

## Background

# Background

## Differential Privacy

*Definition 3.1 (( $\epsilon, \delta$ )-DP [7, 8]).* A randomized mechanism  $\mathcal{M}$  satisfies ( $\epsilon, \delta$ )-differential privacy if for every event  $S \subseteq \text{range}(\mathcal{M})$  and for all  $D \sim D' \in \mathcal{D}^n$ ,

$$\Pr[\mathcal{M}(D) \in S] \leq \exp(\epsilon) \Pr[\mathcal{M}(D') \in S] + \delta.$$

- When  $\delta=0$ , said to be pure-DP

*Definition 3.2 ( $L_1$  and  $L_2$  sensitivity).* Let  $q : \mathcal{D}^n \rightarrow \mathbb{R}^d$  be a query function. The  $L_1$  (resp.  $L_2$ ) sensitivity of  $q$ , denoted by  $\Delta_1(q)$  (resp.,  $\Delta_2(q)$ ) is defined as

# Background

## Gaussian Mechanism

**THEOREM 3.3 (GAUSSIAN MECHANISM [9]).** Let  $\epsilon \in (0, 1)$  be arbitrary and  $q$  be a query function with  $L_2$  sensitivity of  $\Delta_2(q)$ . The Gaussian Mechanism, which returns  $q(D) + N(0, \sigma^2)$ , with

$$\sigma \geq \frac{\Delta_2(q)}{\epsilon} \sqrt{2 \ln(1.25/\delta)} \quad (3)$$

is ( $\epsilon, \delta$ )-differentially private.

# Background

## Concentrated DP

$\rho$ -zCDP,

- for output  $o$  in  $\text{range}(\mathcal{M})$ , the privacy loss random variable  $Z$  of  $\mathcal{M}$  is

$$Z = \log \frac{\Pr[\mathcal{M}(D) = o]}{\Pr[\mathcal{M}(D') = o]}$$

- $\rho$ -zCDP imposes a bound on the moment generating function of  $Z$  and requires it to be bound around 0
- Formally it satisfies:

$$e^{D_\alpha(\mathcal{M}(D) \parallel \mathcal{M}(D'))} = \mathbb{E} \left[ e^{(\alpha-1)Z} \right] \leq e^{(\alpha-1)\alpha\rho}, \quad \forall \alpha \in (1, \infty)$$

# Background

## Concentrated DP Lemmas

- Composition

**LEMMA 3.5 ([4]).** *Suppose two mechanisms satisfy  $\rho_1$ -zCDP and  $\rho_2$ -zCDP, then their composition satisfies  $(\rho_1 + \rho_2)$ -zCDP.*

# Background

## Concentrated DP Lemmas

- zCDP Gaussian Mechanism

**LEMMA 3.6 ([4]).** *The Gaussian mechanism, which returns  $q(D) + N(0, \sigma^2)$  satisfies  $\Delta_2(q)^2 / (2\sigma^2)$ -zCDP.*



# Background

## Concentrated DP Lemmas

- zCDP - pure-DP conversion

**LEMMA 3.7 ([4]).** *If  $\mathcal{M}$  satisfies  $\epsilon$ -differential privacy, then  $\mathcal{M}$  satisfies  $(\frac{1}{2}\epsilon^2)$ -zCDP.*

# Background

## Concentrated DP Lemmas

- zCDP - (e,d)-DP

# Background

## NoisyMax

---

**Algorithm 1:** NOISYMAX( $\Omega, \Delta_1(f), \epsilon$ )

---

**Input:**  $\Omega$ : a set of candidates,  $\Delta_1(f)$ : sensitivity of  $f$ ,  $\epsilon$ :  
privacy budget for pure differential privacy

1  $\tilde{\Omega} = \{\tilde{v}_i = v + \text{Lap}(\Delta_1(f)/\epsilon) : v \in \Omega, i \in [|\Omega|]\}$

2 **return**  $\arg \max_{j \in [|\Omega|]} \tilde{v}_j$

---

NoisyMin = NoisyMax on  $-f(w)$

# NoisyMax

## Implementation

$\Phi$

- Initialize  $\Phi$  with  $m$  equally-spaced pts b/t 0 and  $a_{\max}$
- Proposed to update  $a_{\max}$  in  $\Phi$  every  $r$  iters

$$\alpha_{\max} = (1 + \eta) \max(\alpha_t, \alpha_{t-1}, \dots, \alpha_{t-\tau+1}),$$

- Set  $m=20, r=10, \eta=0.1, a_{\max}=2$

## The Algorithm

DP-AGD

# The Algorithm

## Three Main Parts:

1. Calculate the Noisy Gradient
2. Determine the best step size
3. Adjust the privacy budget (if needed)

# The Algorithm

- Consider ERM

$$\underset{\mathbf{w} \in \mathcal{C}}{\text{minimize}} \quad f(\mathbf{w}; D) := \frac{1}{n} \sum_{i=1}^n \ell(\mathbf{w}; d_i)$$

# The Algorithm

**Input:** privacy budget  $\rho_{\text{nmax}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}$ , budget increase rate  $\gamma$ , clipping thresholds  $C_{\text{obj}}, C_{\text{grad}}$ , data  $\{d_1, \dots, d_n\}$ , objective function  $f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)$

- $\rho_{\text{nmax}}$ : privacy budget in terms of p-zCDP
- $\rho_{\text{ng}}$ : initial privacy budget for the gradient approximation
- $\epsilon_{\text{tot}}, \delta_{\text{tot}}$ : the budget parameters re (E,delta)-DP
- $\gamma$ : budget increase rate
- $C_{\text{obj}}$  and  $C_{\text{grad}}$ : clipping thresholds for gradient calc and obj fnxn
- $\{d_1, \dots, d_n\}$ : set of  $n$  observations where each  $d_i = (x_i, y_i)$
- objective function: ERM

# The Algorithm

---

## Algorithm 2: DP-AGD

---

**Input:** privacy budget  $\rho_{\text{nmax}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}$ , budget increase rate  $\gamma$ , clipping thresholds  $C_{\text{obj}}, C_{\text{grad}}$ , data  $\{d_1, \dots, d_n\}$ , objective function  $f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)$

- 1 Initialize  $\mathbf{w}_0$  and  $\Phi$

# The Algorithm

---

## Algorithm 2: DP-AGD

---

**Input:** privacy budget  $\rho_{\text{nmax}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}$ , budget increase rate  $\gamma$ , clipping thresholds  $C_{\text{obj}}, C_{\text{grad}}$ , data  $\{d_1, \dots, d_n\}$ , objective function  $f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)$

- 1 Initialize  $\mathbf{w}_0$  and  $\Phi$
- 2  $t \leftarrow 0, \rho \leftarrow \text{solve (5) for } \rho$  // To compare to  $(\epsilon, \delta)$ -DP Algs

$$\epsilon_{\text{tot}} \geq \rho + 2\sqrt{\rho \log(1/\delta_{\text{tot}})}.$$

(Lemma 3.8)

# The Algorithm

---

## Algorithm 2: DP-AGD

---

**Input:** privacy budget  $\rho_{\text{nmax}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}$ , budget increase rate  $\gamma$ , clipping thresholds  $C_{\text{obj}}, C_{\text{grad}}$ , data  $\{d_1, \dots, d_n\}$ , objective function  $f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)$

- 1 Initialize  $\mathbf{w}_0$  and  $\Phi$
- 2  $t \leftarrow 0, \rho \leftarrow$  solve (5) for  $\rho$  // To compare to  $(\epsilon, \delta)$ -DP Algs
- 3 **while**  $\rho > 0$  **do**

- Checking use of privacy budget

### Three Main Parts:

1. Calculate the Noisy Gradient
2. Determine the best step size
3. Adjust the privacy budget (if needed)

# The Algorithm

## 1) Private Gradient Approximation

| –   |  |
|---|--|
| <hr/> <b>Algorithm 2: DP-AGD</b> <hr/> <b>Input:</b> privacy budget $\rho_{\text{nmax}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}$ , budget increase rate $\gamma$ , clipping thresholds $C_{\text{obj}}, C_{\text{grad}}$ , data $\{d_1, \dots, d_n\}$ , objective function $f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)$ <ol style="list-style-type: none"> <li>1 Initialize <math>\mathbf{w}_0</math> and <math>\Phi</math></li> <li>2 <math>t \leftarrow 0, \rho \leftarrow \text{solve (5) for } \rho</math> // To compare to <math>(\epsilon, \delta)</math>-DP Algs</li> <li>3 <b>while</b> <math>\rho &gt; 0</math> <b>do</b></li> </ol> | + Gradient Clipping technique<br>(instead of assuming $\text{norm}(\mathbf{x}) \leq 1$ )<br><br>+ variance of noise depends on |

# The Algorithm

## 1) Private Gradient Approximation

| –   |                 |
|---|-----------------|
| <hr/> <b>Algorithm 2: DP-AGD</b> <hr/> <b>Input:</b> privacy budget $\rho_{\text{nmax}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}$ , budget increase rate $\gamma$ , clipping thresholds $C_{\text{obj}}, C_{\text{grad}}$ , data $\{d_1, \dots, d_n\}$ , objective function $f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)$ <ol style="list-style-type: none"> <li>1 Initialize <math>\mathbf{w}_0</math> and <math>\Phi</math></li> <li>2 <math>t \leftarrow 0, \rho \leftarrow \text{solve (5) for } \rho</math> // To compare to <math>(\epsilon, \delta)</math>-DP Algs</li> <li>3 <b>while</b> <math>\rho &gt; 0</math> <b>do</b></li> <li>4     <math>i \leftarrow 0</math></li> <li>5     <math>\mathbf{g}_t \leftarrow \sum_{i=1}^n \left( \nabla \ell(\mathbf{w}_t; d_i) / \max(1, \frac{\ \nabla \ell(\mathbf{w}_t)\ _2}{C_{\text{grad}}}) \right)</math></li> <li>6     <math>\tilde{\mathbf{g}}_t \leftarrow \mathbf{g}_t + N(0, (C_{\text{grad}}^2 / 2\rho_{\text{ng}})\mathbf{I})</math></li> </ol> | + Add the noise |



# The Algorithm

## 1) Private Gradient Approximation

|  |                            |
|--|----------------------------|
| -  | -                          |
| <hr/> <p><b>Algorithm 2: DP-AGD</b></p> <hr/> <p><b>Input:</b> privacy budget <math>\rho_{\text{nmax}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}</math>, budget increase rate <math>\gamma</math>, clipping thresholds <math>C_{\text{obj}}, C_{\text{grad}}</math>, data <math>\{d_1, \dots, d_n\}</math>, objective function <math>f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)</math></p> <ol style="list-style-type: none"> <li>1 Initialize <math>\mathbf{w}_0</math> and <math>\Phi</math></li> <li>2 <math>t \leftarrow 0, \rho \leftarrow \text{solve (5) for } \rho</math> // To compare to <math>(\epsilon, \delta)</math>-DP Algs</li> <li>3 <b>while</b> <math>\rho &gt; 0</math> <b>do</b></li> <li>4     <math>i \leftarrow 0</math></li> <li>5     <math>\mathbf{g}_t \leftarrow \sum_{i=1}^n \left( \nabla \ell(\mathbf{w}_t; d_i) / \max(1, \frac{\ \nabla \ell(\mathbf{w}_t)\ _2}{C_{\text{grad}}}) \right)</math></li> <li>6     <math>\tilde{\mathbf{g}}_t \leftarrow \mathbf{g}_t + N(0, (C_{\text{grad}}^2 / 2\rho_{\text{ng}})\mathbf{I})</math></li> <li>7     <math>\rho \leftarrow \rho - \rho_{\text{ng}}</math></li> </ol> | <p>+ Adjust the budget</p> |

# The Algorithm

## 1) Private Gradient Approximation

|  |  |
|--|--|
|  |  |
|--|--|

# The Algorithm

## 2) Step Size Selection

|  |  |
|--|--|
|  |  |
| <hr/> <p><b>Algorithm 2:</b> DP-AGD</p> <hr/> <p><b>Input:</b> privacy budget <math>\rho_{\text{nmax}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}</math>, budget increase rate <math>\gamma</math>, clipping thresholds <math>C_{\text{obj}}, C_{\text{grad}}</math>, data <math>\{d_1, \dots, d_n\}</math>, objective function <math>f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)</math></p> <p>1 Initialize <math>\mathbf{w}_0</math> and <math>\Phi</math></p> <p>2 <math>t \leftarrow 0, \rho \leftarrow \text{solve (5) for } \rho</math> // To compare to <math>(\epsilon, \delta)</math>-DP Algs</p> <p>3 <b>while</b> <math>\rho &gt; 0</math> <b>do</b></p> <p>4   <math>i \leftarrow 0</math></p> <p>5   <math>\mathbf{g}_t \leftarrow \sum_{i=1}^n \left( \nabla \ell(\mathbf{w}_t; d_i) / \max(1, \frac{\ \nabla \ell(\mathbf{w}_t)\ _2}{C_{\text{grad}}}) \right)</math></p> <p>6   <math>\tilde{\mathbf{g}}_t \leftarrow \mathbf{g}_t + N(0, (C_{\text{grad}}^2 / 2\rho_{\text{ng}})\mathbf{I})</math></p> <p>7   <math>\rho \leftarrow \rho - \rho_{\text{ng}}</math></p> <p>8   <math>\tilde{\mathbf{g}}_t \leftarrow \tilde{\mathbf{g}}_t / \ \tilde{\mathbf{g}}_t\ _2</math></p> <p>9   <b>while</b> <math>i = 0</math> <b>do</b></p> <p>10     <math>\Omega = \{f(\mathbf{w}_t - \alpha \tilde{\mathbf{g}}_t) : \alpha \in \Phi\}</math></p> <p>11     <math>\rho \leftarrow \rho - \rho_{\text{nmax}}</math></p> <p>12     <math>i \leftarrow \text{NOISYMAX}(-\Omega, C_{\text{obj}}, \sqrt{2\rho_{\text{nmax}}})</math></p> | <p>+ Find optimal step size</p> <p>+ NoisyMax(set of candidates, sensitivity, privacy budget for pure DP - Lemma 3.7)</p> <p>+ Control the variance of the updates</p> |

# The Algorithm

## 2) Step Size Selection

| —  |  |
|--|--|
| <p><b>Algorithm 2:</b> DP-AGD</p> <p><b>Input:</b> privacy budget <math>\rho_{\text{max}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}</math>, budget increase rate <math>\gamma</math>, clipping thresholds <math>C_{\text{obj}}, C_{\text{grad}}</math>, data <math>\{d_1, \dots, d_n\}</math>, objective function <math>f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)</math></p> <ol style="list-style-type: none"> <li>1 Initialize <math>\mathbf{w}_0</math> and <math>\Phi</math></li> <li>2 <math>t \leftarrow 0, \rho \leftarrow \text{solve (5) for } \rho</math> // To compare to <math>(\epsilon, \delta)</math>-DP Algs</li> <li>3 <b>while</b> <math>\rho &gt; 0</math> <b>do</b></li> <li>4   <math>i \leftarrow 0</math></li> <li>5   <math>\mathbf{g}_t \leftarrow \sum_{i=1}^n \left( \nabla \ell(\mathbf{w}_t; d_i) / \max(1, \frac{\ \nabla \ell(\mathbf{w}_t)\ _2}{C_{\text{grad}}}) \right)</math></li> <li>6   <math>\tilde{\mathbf{g}}_t \leftarrow \mathbf{g}_t + N(0, (C_{\text{grad}}^2 / 2\rho_{\text{ng}})\mathbf{I})</math></li> <li>7   <math>\rho \leftarrow \rho - \rho_{\text{ng}}</math></li> <li>8   <math>\tilde{\mathbf{g}}_t \leftarrow \tilde{\mathbf{g}}_t / \ \tilde{\mathbf{g}}_t\ _2</math></li> <li>9   <b>while</b> <math>i = 0</math> <b>do</b></li> </ol> | <p>Step 14: Check use of privacy budget before</p> |

# The Algorithm

## 3) Adoptive Noise Reduction - Gradient Averaging

| —  |  |
|--|--|
| <ol style="list-style-type: none"> <li>15   <b>else</b></li> <li>16       <math>\rho_{\text{old}} \leftarrow \rho_{\text{ng}}</math></li> <li>17       <math>\rho_{\text{ng}} \leftarrow (1 + \gamma)\rho_{\text{ng}}</math></li> <li>18       <math>\tilde{\mathbf{g}}_t \leftarrow \text{GRADAVG}(\rho_{\text{old}}, \rho_{\text{ng}}, \mathbf{g}_t, \tilde{\mathbf{g}}_t, C_{\text{grad}})</math></li> <li>19       <math>\rho \leftarrow \rho - (\rho_{\text{ng}} - \rho_{\text{old}})</math></li> </ol><br><ol style="list-style-type: none"> <li>25 <b>Function</b> GRADAVG(<math>\rho_{\text{old}}, \rho_H, \mathbf{g}, \tilde{\mathbf{g}}, C_{\text{grad}}</math>):</li> <li>26   <math>\tilde{\mathbf{g}}_2 \leftarrow \mathbf{g} + N(0, (\frac{C_{\text{grad}}^2}{2(\rho_H - \rho_{\text{old}})})\mathbf{I})</math></li> <li>27   <math>\tilde{\mathbf{S}} \leftarrow \frac{\rho_{\text{old}}\tilde{\mathbf{g}} + (\rho_H - \rho_{\text{old}})\tilde{\mathbf{g}}_2}{\rho_H}</math></li> <li>28   <b>return</b> <math>\tilde{\mathbf{S}}</math></li> <li>29 <b>end</b></li> </ol> | <p>+ Increase <math>p_t</math> to <math>p_{t+1}</math></p> <p>+ Only use <math>p_{t+1} - p_t</math></p> <p>+ Average the gradient from steps <math>t</math> and <math>t+1</math> to only use a total of <math>p_{t+1}</math> across both</p> |

# The Algorithm

| —  | —   |
|--|---|
| <hr/> <b>Algorithm 2:</b> DP-AGD <hr/> <p><b>Input:</b> privacy budget <math>\rho_{\text{nmax}}, \rho_{\text{ng}}, \epsilon_{\text{tot}}, \delta_{\text{tot}}</math>, budget increase rate <math>\gamma</math>, clipping thresholds <math>C_{\text{obj}}, C_{\text{grad}}</math>, data <math>\{d_1, \dots, d_n\}</math>, objective function <math>f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)</math></p> <pre> 1 Initialize <math>\mathbf{w}_0</math> and <math>\Phi</math> 2 <math>t \leftarrow 0, \rho \leftarrow \text{solve (5) for } \rho</math> // To compare to <math>(\epsilon, \delta)</math>-DP Algs 3 <b>while</b> <math>\rho &gt; 0</math> <b>do</b> 4   <math>i \leftarrow 0</math> 5   <math>\mathbf{g}_t \leftarrow \sum_{i=1}^n \left( \nabla \ell(\mathbf{w}_t; d_i) / \max(1, \frac{\ \nabla \ell(\mathbf{w}_t)\ _2}{C_{\text{grad}}}) \right)</math> 6   <math>\tilde{\mathbf{g}}_t \leftarrow \mathbf{g}_t + N(0, (C_{\text{grad}}^2 / 2\rho_{\text{ng}})\mathbf{I})</math> 7   <math>\rho \leftarrow \rho - \rho_{\text{ng}}</math> 8   <math>\tilde{\mathbf{g}}_t \leftarrow \tilde{\mathbf{g}}_t / \ \tilde{\mathbf{g}}_t\ _2</math> 9   <b>while</b> <math>i = 0</math> <b>do</b> 10    <math>\Omega = \{f(\mathbf{w}_t - \alpha \tilde{\mathbf{g}}_t) : \alpha \in \Phi\}</math> 11    <math>\rho \leftarrow \rho - \rho_{\text{nmax}}</math> 12    <math>i \leftarrow \text{NOISYMAX}(-\Omega, C_{\text{obj}}, \sqrt{2\rho_{\text{nmax}}})</math> 13    <b>if</b> <math>i &gt; 0</math> <b>then</b> 14      <b>if</b> <math>\rho &gt; 0</math> <b>then</b> <math>\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t - \alpha_i \tilde{\mathbf{g}}_t</math> 15    <b>else</b> 16      <math>\rho_{\text{old}} \leftarrow \rho_{\text{ng}}</math> 17      <math>\rho_{\text{ng}} \leftarrow (1 + \gamma)\rho_{\text{ng}}</math> 18      <math>\tilde{\mathbf{g}}_t \leftarrow \text{GRADAVG}(\rho_{\text{old}}, \rho_{\text{ng}}, \mathbf{g}_t, \tilde{\mathbf{g}}_t, C_{\text{grad}})</math> 19      <math>\rho \leftarrow \rho - (\rho_{\text{ng}} - \rho_{\text{old}})</math> 20    <b>end</b> 21  <b>end</b> 22  <math>t \leftarrow t + 1</math> </pre> | <p>+ loop, dynamically computing budget <math>\rho</math> (lines 19, 11, 7)</p> <p>+ thanks to composition properties</p> |

# The Algorithm

| -   | - |
|---|---|
| <hr/> <b>Algorithm 2:</b> DP-AGD <hr/> <b>Input:</b> privacy budget $\rho_{nmax}, \rho_{ng}, \epsilon_{tot}, \delta_{tot}$ , budget increase rate $\gamma$ , clipping thresholds $C_{obj}, C_{grad}$ , data $\{d_1, \dots, d_n\}$ , objective function $f(\mathbf{w}) = \sum_{i=1}^n \ell(\mathbf{w}; d_i)$ |   |

```
def agd_rho(X, y, rho, eps_total, delta, grad_func, loss_func, test_func,
            obj_clip, grad_clip, reg_coeff=0.0, batch_size=-1, exp_dec=1.0,
            gamma=0.1, splits=60, verbose=False):
    N, dim = X.shape

    # parameters
    eps_nmax = (eps_total * 0.5) / splits
    sigma = compute_sigma(eps_nmax, delta, grad_clip)

    # intial privacy budget
    rho_nmax = 0.5 * (eps_nmax**2)
    rho_ng = (grad_clip**2) / (2.0 * sigma**2)
    # rho_ng = rho_nmax

    # initialize the parameter vector
    w = np.zeros(dim)
    t = 0
    chosen_step_sizes = []
    max_step_size = 2.0
    n_candidate = 20
```

```

def agd_rho():
    ...
    while rho > 0:
        if verbose:
            loss = loss_func(w, X, y) / N
            acc = test_func(w, X, y)
            print "[{}] loss: {:.5f} acc: {:.2f}".format(t, loss, acc*100)
        if batch_size > 0:
            # build a mini-batch
            rand_idx = np.random.choice(N, size=batch_size, replace=False)
            mini_X = X[rand_idx, :]
            mini_y = y[rand_idx]
        else:
            mini_X = X
            mini_y = y

        # non-private (clipped) gradient
        grad = grad_func(w, mini_X, mini_y, grad_clip)

        sigma = grad_clip / math.sqrt(2.0 * rho_ng)
        noisy_grad = grad + sigma * np.random.randn(dim)
        noisy_unnorm = np.copy(noisy_grad)
        noisy_grad /= np.linalg.norm(noisy_grad)

        rho -= rho_ng

```

```

def agd_rho():
    ...
    while rho > 0:
        ...
        while idx == 0:
            # test if this is a descent direction
            step_sizes = np.linspace(0, max_step_size, n_candidate+1)
            candidate = [w - step * noisy_grad for step in step_sizes]
            scores = [loss_func(theta, mini_X, mini_y, clip=obj_clip)
                      for theta in candidate]
            scores[0] *= exp_dec

            # deduct the privacy budget used for noisy max
            lmbda = obj_clip/math.sqrt(2.0 * rho_nmax)
            idx, _ = noisy_max(scores, lmbda, bmin=True)
            rho -= rho_nmax

```

```

def agd_rho():
    ...
    while rho > 0:
        ...
        while idx == 0:
            ...
            # used up the budget
            if rho < 0:
                break

            if idx > 0:
                # don't do the update when the remain budget is insufficient
                if rho >= 0:
                    w[:] = candidate[idx-1]
                    rho -= rho_ng
            else:
                rho_old = rho_ng
                rho_ng *= (1.0 + gamma)
                # max_step_size *= 0.9
                if verbose:
                    print "\tbad gradient: resample (rho_ng={})".format(rho_ng)

                noisy_grad = grad_avg(rho_old, rho_ng, grad, noisy_unnorm,
                                       grad_clip)
                noisy_grad /= np.linalg.norm(noisy_grad)
                rho -= (rho_ng - rho_old)

            if reg_coeff > 0:
                noisy_grad += reg_coeff * w

```

```

def agd_rho():
    ...
    while rho > 0:
        ...
        while idx == 0:
            ...
            chosen_step_sizes.append(step_sizes[idx])
            t += 1

            if (t % 10) == 0:
                max_step_size = min(1.1*max(chosen_step_sizes), 2.0)
                del chosen_step_sizes[:]

def grad_avg(rho_old, rho_H, true_grad, noisy_grad, grad_clip):
    dim = true_grad.shape[0]
    sigma = grad_clip / math.sqrt(2.0 * (rho_H - rho_old))

    # new estimate
    g_2 = true_grad + sigma * np.random.randn(dim)

    beta = rho_old / rho_H
    # weighted average
    s_tilde = beta * noisy_grad + (1.0 - beta) * g_2

    return s_tilde

```



# Experimental Results

# Experimental Results

## Experimental Results

### Baselines:

**ObjPert:** Objective Perturbation, add linear perturbation term to the objective function and solve using non-private solver

**OutPert:** Output Perturbation, perturb the output after batch GD for predetermined steps

**PrivGene:** DP genetic algorithm

**SGD-Adv:** DP algorithm using advanced composition with privacy amplification output (upper bound)

**SGD-MA:** Uses an improved composition method called moments accountant

**NonPrivate:** Non-private L-BFGS algorithm

**Majority:** Chooses majority class

- All known DP algorithms use a predetermined privacy budget sequence

# Experimental Results

## Experiment Design

- Report the classification accuracy and final objective value
- Perform 20 repeats of 5-fold CV

### Preprocessing:

- one-hot encoding of all binary variables
- rescale numerical variables on  $[0,1]$
- for *ChIPert*: normalize to unit norm (per requirements)

# Experimental Results

## Effects of Parameters

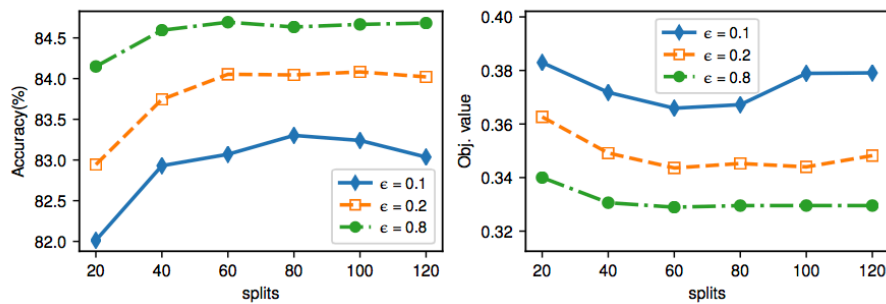
- Ran DP-AGD on Adult dataset
- Overall: robust to parameters - moderate settings have small effect

# Experimental Results

## Effects of Parameters (DP-AGD on Adult)

- Splits (T) : Estimate initial privacy budget via number of splits

$$\epsilon_{\text{ng}} = \epsilon_{\text{ng}} = \frac{\epsilon_{\text{tot}}}{2 \cdot \text{splits}}$$



(a) Effect of splits (left: accuracy, right: obj. value)

- Less effect when budget is large
- Too small & Too large

# Experimental Results

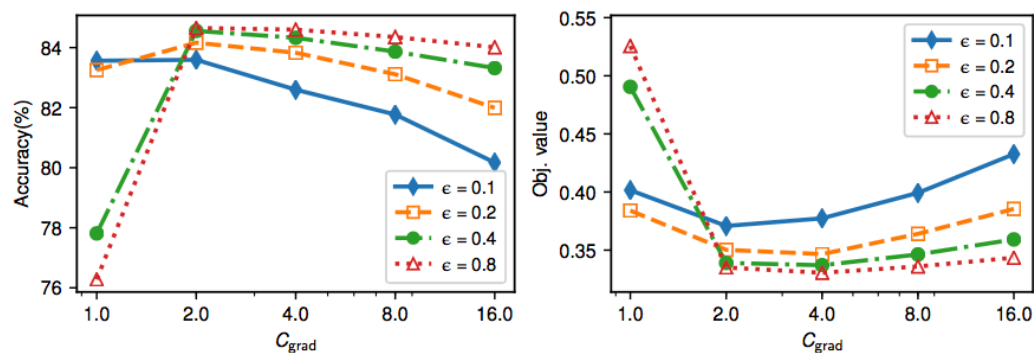
## Effects of Parameters (DP-AGD on Adult)

•  $C_{\text{grad}}$

# Experimental Results

## Effects of Parameters (DP-AGD on Adult)

- $C_{\text{grad}}$



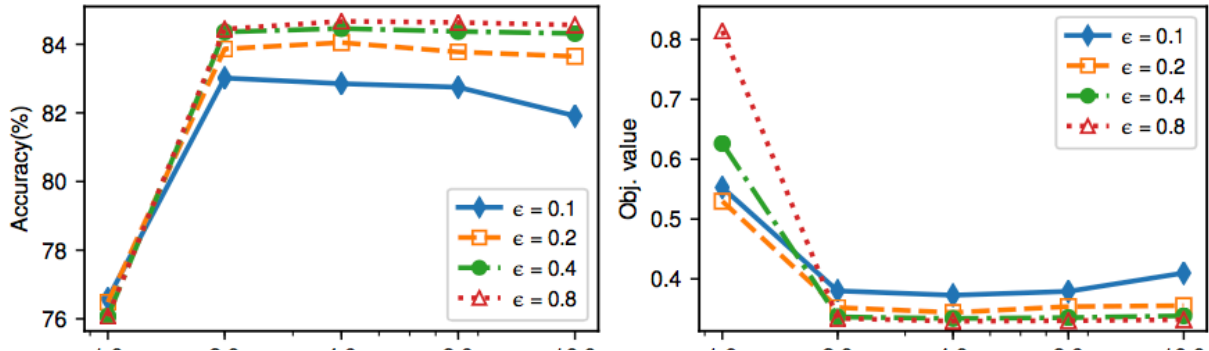
**(c) Effect of  $C_{\text{grad}}$  (left: accuracy, right: obj. value)**

- $C_x$  are the clipping parameters : sets a bound on the sensitivity to satisfy zCDP
- $C_{\text{grad}}$  refers to the gradient calculation
- Too small: low sensitivity, but can cause too much information loss
- Too large: high sensitivity, results in too much noise

# Experimental Results

## Effects of Parameters (DP-AGD on Adult)

- $C_{obj}$



# Experimental Results

## Logistic Regression

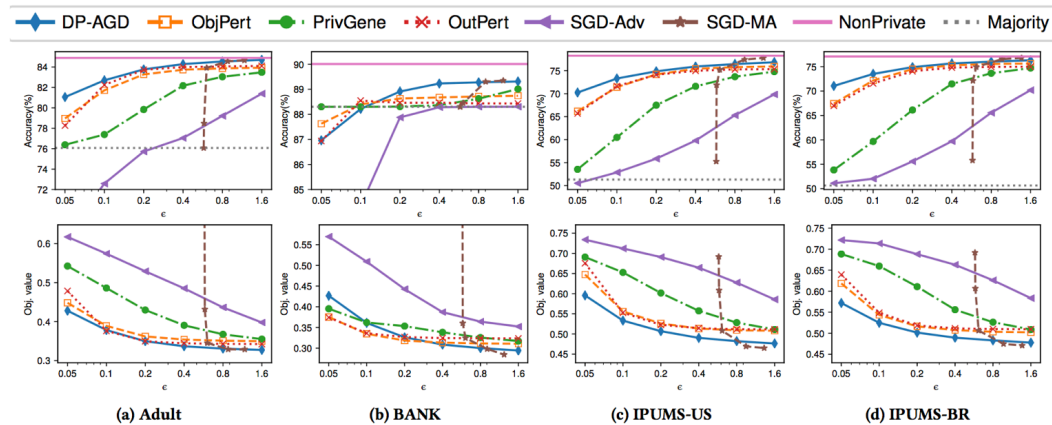


Figure 2: Logistic regression by varying  $\epsilon$  (Top: classification accuracies, Bottom: objective values)

|   |  |
|---|--|
| —   | —  |
| + DP-AGD outperforms / performs competitively<br>+ especially when budget is very small | + SGD-Adv and SGD-Ma depend on T (predetermined)<br>+ SGD-Ma outperforms all others when budget > 0.8 (more splits due to tight bound on privacy loss)<br>--> impractical under strict privacy control |

# Experimental Results

## SVM

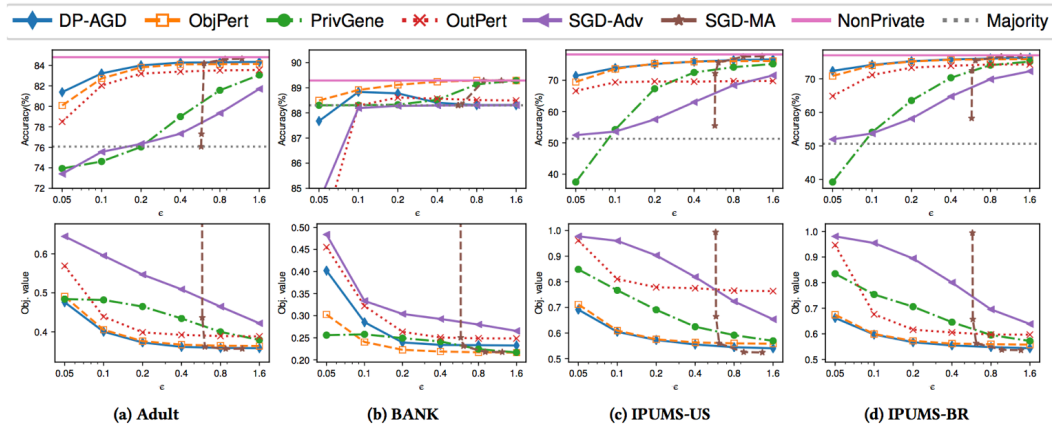
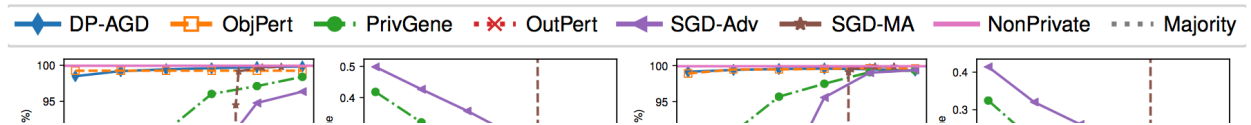


Figure 3: SVM by varying  $\epsilon$  (Top: classification accuracies, Bottom: objective values)

- DP-AGD scores competitive accuracies
  - shows unstable behavior on BANK (acc. decreases despite increased use of budget)
  - Still, objective value consistently decreases as the budget is increased

# Experimental Results

## KDDCup99



## Conclusions:

- First known iterative optimization algorithm for differential privacy
- satisfies zCDP, which is weaker than (e)-DP, but stronger than (e,d)-DP
- Per-iteration privacy budget is adaptively determined during runtime based on utility of privacy-preserving statistics
- Uses gradient averaging to utilize non-descending gradients and preserve the budget
- Outperforms or performs competitively with baselines, with an advantage under strict privacy settings
- Code available: <https://github.com/ppmlguy/DP-AGD> (<https://github.com/ppmlguy/DP-AGD>).
- Citation:

```
@inproceedings{Lee2018ConcentratedDP,
  title={Concentrated Differentially Private Gradient Descent with Adaptive per-Iteration Privacy Budget},
  author={Jaewoo Lee and Daniel Kifer},
  booktitle={KDD},
  year={2018}
}
```

- Lee, J., & Kifer, D. (2018). Concentrated Differentially Private Gradient Descent with Adaptive per-Iteration Privacy Budget. KDD.



# Questions?